

```
In [262]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import RandomizedSearchCV
```

```
In [263]: filepath = "/content/drive/MyDrive/synthetic_ecommerce_data.csv"
```

```
In [264]: df = pd.read_csv(filepath, parse_dates=['Transaction_Date'])
```

```
##EDA
```

| Out[265]: | Transaction_ID | Customer_ID | Product_ID | Transaction_Date |
|-----------|--------------------------------------|----------------|-------------|------------------|
| 0 | 8b460852-7c64-46fa-998b-b0976879d082 | Customer_65 | Product_224 | 2024-10-06 |
| 1 | 418612e7-8744-4ba3-bb0c-105b47e2a968 | Customer_1910 | Product_584 | 2024-10-29 |
| 2 | 5bc3b98f-cb0c-4b12-947c-df8bbb35a73e | Customer_2306 | Product_374 | 2024-04-04 |
| 3 | 28fb67c8-e8c0-447a-841c-f760730de0eb | Customer_17206 | Product_220 | 2024-08-25 |
| 4 | 8bee087a-a8a9-45bb-89d7-04d1710f1b00 | Customer_16033 | Product_358 | 2024-05-05 |

```
In [266]: df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Transaction_ID    100000 non-null   object  
 1   Customer_ID       100000 non-null   object  
 2   Product_ID        100000 non-null   object  
 3   Transaction_Date  100000 non-null   datetime64[ns]
 4   Units_Sold        100000 non-null   int64   
 5   Discount_Applied  100000 non-null   float64 
 6   Revenue            100000 non-null   float64 
 7   Clicks             100000 non-null   int64   
 8   Impressions        100000 non-null   int64   
 9   Conversion_Rate   100000 non-null   float64 
 10  Category           100000 non-null   object  
 11  Region             100000 non-null   object  
 12  Ad_CTR             100000 non-null   float64 
 13  Ad_CPC             100000 non-null   float64 
 14  Ad_Spend            100000 non-null   float64 
dtypes: datetime64[ns](1), float64(6), int64(3), object(5)
memory usage: 11.4+ MB
```

```
In [267]: # nothing's missing lowkey, some cols are objects (categories to handle before training on)
```

```
In [268]: df.isnull().sum()  
# confirmed
```

| Out[268]: | |
|------------------|---|
| | 0 |
| Transaction_ID | 0 |
| Customer_ID | 0 |
| Product_ID | 0 |
| Transaction_Date | 0 |
| Units_Sold | 0 |
| Discount_Applied | 0 |
| Revenue | 0 |
| Clicks | 0 |
| Impressions | 0 |
| Conversion_Rate | 0 |
| Category | 0 |
| Region | 0 |
| AdCTR | 0 |
| AdCPC | 0 |
| AdSpend | 0 |

dtype: int64

Outlier detection


if they're outliers or justifiable data points

```
In [269]: plt.figure(figsize=(12, 5))

# Plotting Ad Spend and Revenue outliers
plt.subplot(1, 2, 1)
sns.boxplot(y=df['Ad_Spend'], color='skyblue')
plt.title('Outliers in Ad Spend')

plt.subplot(1, 2, 2)
sns.boxplot(y=df['Revenue'], color='salmon')
plt.title('Outliers in Revenue')

plt.tight_layout()
plt.show()
```



Several high-revenue data points are "outliers", think of a model that's robust to this situation

Feature engineering

```
In [270]: df['Month'] = df['Transaction_Date'].dt.month
df['DayOfWeek'] = df['Transaction_Date'].dt.dayofweek

In [271]: # engineering the campaigns out of date, cat, region
data = df.groupby(['Transaction_Date', 'Category', 'Region', 'Month',
'DayOfWeek']).agg({
    'Ad_Spend': 'sum',
    'Clicks': 'sum',
    'Impressions': 'sum',
    'Revenue': 'sum',
    'Units_Sold': 'sum',
    'Ad_CTR': 'sum',
    'Ad_CPC': 'sum'
}).reset_index()
```

```
In [272]: data['Spend_Efficiency'] = data['Ad_Spend'] * data['Clicks']
```

```
In # features and y : target column  
[273]: X = data[['Ad_Spend', 'Clicks', 'Impressions', 'Month', 'DayOfWeek',  
'Category', 'Region', 'Spend_Efficiency', 'AdCTR', 'AdCPC']]  
y = data['Revenue']
```

```
In X  
[274]:
```

| Out[274]: | | Ad_Spend | Clicks | Impressions | Month | DayOfWeek | Category |
|-----------|------|----------|--------|-------------|-------|-----------|-----------------|
| | 0 | 1837.26 | 389 | 3365 | 12 | 3 | Books |
| | 1 | 2385.95 | 405 | 4477 | 12 | 3 | Books |
| | 2 | 2910.23 | 517 | 6199 | 12 | 3 | Books |
| | 3 | 2498.42 | 403 | 4692 | 12 | 3 | Clothing |
| | 4 | 1845.82 | 309 | 4542 | 12 | 3 | Clothing |
| | ... | ... | ... | ... | ... | ... | ... |
| | 5485 | 2093.20 | 256 | 3343 | 12 | 4 | Home Appliances |
| | 5486 | 2224.22 | 409 | 4165 | 12 | 4 | Home Appliances |
| | 5487 | 1797.73 | 430 | 5375 | 12 | 4 | Toys |
| | 5488 | 1581.18 | 409 | 3015 | 12 | 4 | Toys |
| | 5489 | 1087.23 | 339 | 2919 | 12 | 4 | Toys |

5490 rows × 10 columns

```
In # Encoding categories (transforming the Xs that are categories to  
[275]: numericals)  
X = pd.get_dummies(X, columns=['Category', 'Region'], drop_first=True)  
# X : as expected
```

```
In X_train_full, X_test, y_train_full, y_test = train_test_split(X, y,  
[276]: test_size=0.2, random_state=42)  
X_train, X_val, y_train, y_val = train_test_split(X_train_full,  
y_train_full, test_size=0.2, random_state=42)  
  
print(f"Train size: {X_train.shape[0]}, Val size: {X_val.shape[0]}, Test  
size: {X_test.shape[0]}")
```

Train size: 3513, Val size: 879, Test size: 1098

Further preprocessing

normalization, standardization

```
In [277]: num_features = ['Ad_Spend', 'Clicks', 'Impressions', 'Month',
'DayOfWeek']

# scaler at work
scaler = StandardScaler()
# init the new X_train , y_train
X_train_scaled = X_train.copy()
X_val_scaled = X_val.copy()
X_test_scaled = X_test.copy()


X_train_scaled[num_features] =
scaler.fit_transform(X_train[num_features])
X_val_scaled[num_features] = scaler.transform(X_val[num_features])
X_test_scaled[num_features] = scaler.transform(X_test[num_features])
```

Marketing dynamics analysis


```
In # correlation ? between category budget and revenue  
[278]: for cat in data['Category'].unique():  
    subset = data[data['Category'] == cat]  
    correlation = subset['Ad_Spend'].corr(subset['Revenue'])  
    print(f"Correlation for {cat}: {correlation:.2f}")  
  
plt.figure(figsize=(10, 6))  
sns.lmplot(x='Ad_Spend', y='Revenue', hue='Category', data=data,  
            aspect=1.5, scatter_kws={'alpha':0.3})  
plt.title('Revenue Response per Dollar Spent by Category')  
plt.show()
```

```
Correlation for Books: 0.59  
Correlation for Clothing: 0.58  
Correlation for Electronics: 0.58  
Correlation for Home Appliances: 0.56  
Correlation for Toys: 0.63
```

<Figure size 1000x600 with 0 Axes>



```
In # List of numeric features to test  
[279]: features_to_plot = ['Ad_Spend', 'Clicks', 'Impressions', 'Ad_CTR',  
'Ad_CPC']  
  
fig, axes = plt.subplots(nrows=2, ncols=3, figsize=(18, 10))  
axes = axes.flatten()  
  
for i, col in enumerate(features_to_plot):  
    sns.regplot(data=df, x=col, y='Revenue', ax=axes[i],  
                scatter_kws={'alpha':0.3}, line_kws={'color':'red'})  
    axes[i].set_title(f'{col} vs Revenue')  
  
plt.tight_layout()  
plt.show()
```




```
In # df.dtypes  
[280]:
```

Global correlations

```
In [281]: # cols_excluded_heat_map = ['Transaction_ID', 'Customer_ID',
#                                'Transaction_Date', 'Product_ID']
# only numerics are left
data_HM = pd.concat([X, y], axis=1)

plt.figure(figsize=(12, 8))
sns.heatmap(pd.concat([X, y], axis=1).corr(), annot=True,
            cmap='coolwarm', fmt=".2f", linewidths=0.5)
plt.title("Correlation Matrix: Marketing Metrics & Revenue")
plt.show()
```



##Model training

```
In [284]: from xgboost import XGBRegressor
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import BaggingRegressor, RandomForestRegressor

# grid params to try
param_grid = {
    'n_estimators': [100, 300, 500],
    'learning_rate': [0.01, 0.05, 0.1],
    'max_depth': [3, 7, 12],
    'subsample': [0.8]
}

# Grid Search on XGBoost
grid_xgb = GridSearchCV(XGBRegressor(random_state=42), param_grid, cv=4,
scoring='r2')
grid_xgb.fit(X_train_scaled, y_train)

best_xgb = grid_xgb.best_estimator_
best_xgb.score(X_val_scaled, y_val)
```


Out[284]: 0.5304849480000307

```
In [285]: bag_rf = BaggingRegressor(
    estimator=RandomForestRegressor(n_estimators=100),
    n_estimators=10,
    random_state=42
)

# Fit the model on the same scaled training data used for XGBoost
bag_rf.fit(X_train_scaled, y_train)

bag_test_preds = bag_rf.predict(X_test_scaled)
```

```
In [286]:  
import matplotlib.pyplot as plt  
  
importances = best_xgb.feature_importances_  
feat_importances = pd.Series(importances, index=X.columns)  
  
plt.figure(figsize=(10,6))  
feat_importances.nlargest(10).plot(kind='barh', color='skyblue')  
plt.title("Feature importance (XGBoost)")  
plt.xlabel("Importance score")  
plt.show()
```



```
In [287]:  
from sklearn.metrics import r2_score, mean_absolute_error,  
mean_squared_error  
  
# Final prediction on the tst set (which the model NEVER saw)  
final_test_preds = best_xgb.predict(X_test_scaled)  
  
test_r2 = r2_score(y_test, final_test_preds)  
test_mae = mean_absolute_error(y_test, final_test_preds)  
test_rmse = np.sqrt(mean_squared_error(y_test, final_test_preds))  
  
print(f"--- FINAL TEST SET RESULTS ---")  
print(f"R2 Score: {test_r2:.4f}")  
print(f"MAE: ${test_mae:.2f}")  
print(f"RMSE: ${test_rmse:.2f}")
```

--- FINAL TEST SET RESULTS ---

R2 Score: 0.5244

MAE: \$1634.29

RMSE: \$2273.93

```
In [288]: bag_test_r2 = r2_score(y_test, bag_test_preds)
          bag_test_mae = mean_absolute_error(y_test, bag_test_preds)
          bag_test_rmse = np.sqrt(mean_squared_error(y_test, bag_test_preds))

          print(f"--- BAGGING RF TEST SET RESULTS ---")
          print(f"R2 Score: {bag_test_r2:.4f}")
          print(f"MAE: ${bag_test_mae:.2f}")
          print(f"RMSE: ${bag_test_rmse:.2f}")
```

```
--- BAGGING RF TEST SET RESULTS ---
R2 Score: 0.5037
MAE: $1678.93
RMSE: $2322.93
```

XGBoost is better

```
In [290]: import numpy as np

# Train on log values
y_train_log = np.log1p(y_train)
best_xgb.fit(X_train_scaled, y_train_log)

# Predict and transform back to dollars
log_preds = best_xgb.predict(X_test_scaled)
final_preds = np.expm1(log_preds) # inverser log using exp

print(f"New R2: {r2_score(y_test, final_preds)}")
print(f"MAE: ${mean_absolute_error(y_test, final_test_preds)}")
```

```
New R2: 0.5103692168572471
MAE: $1634.2862791901468
```

Changed nothing significant

```
In [291]: from sklearn.metrics import mean_absolute_error,
mean_absolute_percentage_error

# core Eval metrics
mae = mean_absolute_error(y_test, final_test_preds)
mape = mean_absolute_percentage_error(y_test, final_test_preds)

# Calculate 'accuracy' as 1 - MAPE
accuracy_pct = (1 - mape) * 100

# Calculate the 'Relative Error' against the Mean
mean_revenue = y_test.mean()
relative_error = (mae / mean_revenue) * 100

print(f"--- Objective Performance Metrics ---")
print(f"Mean Revenue in Test Set: ${mean_revenue:.2f}")
print(f"Mean Absolute Error (MAE): ${mae:.2f}")
print(f"Mean Absolute Percentage Error (MAPE): {mape*100:.2f}%")
print(f"--- Report Ready Statistics ---")
print(f"Predictive Accuracy (1-MAPE): {accuracy_pct:.2f}%")
print(f"Error as % of Mean Revenue: {relative_error:.2f}%")
```

```
--- Objective Performance Metrics ---
Mean Revenue in Test Set: $10161.77
Mean Absolute Error (MAE): $1634.29
Mean Absolute Percentage Error (MAPE): 17.64%
--- Report Ready Statistics ---
Predictive Accuracy (1-MAPE): 82.36%
Error as % of Mean Revenue: 16.08%
```

```
In [292]: from sklearn.model_selection import learning_curve
from sklearn.metrics import make_scorer, mean_squared_error
import matplotlib.pyplot as plt
import numpy as np

def plot_comprehensive_learning_curves(models, X, y):
    # Define RMSE scorer since it's not a default string option
    rmse_scorer = make_scorer(lambda y_true, y_pred:
        np.sqrt(mean_squared_error(y_true, y_pred)), greater_is_better=False)

    metrics = [
        ('r2', 'R2 Score', False),
        ('neg_mean_absolute_error', 'MAE ($)', True),
        (rmse_scorer, 'RMSE ($)', True)
    ]

    fig, axes = plt.subplots(len(models), len(metrics), figsize=(20, 12))


    for i, (name, model) in enumerate(models.items()):
        for j, (scorer, label, is_error_metric) in enumerate(metrics):
            train_sizes, train_scores, test_scores = learning_curve(
                model, X, y, cv=3, n_jobs=-1,
                train_sizes=np.linspace(0.1, 1.0, 5), scoring=scorer
            )

            # Convert negative scores back to positive for MAE and RMSE
            if is_error_metric:
                train_mean = -np.mean(train_scores, axis=1)
                test_mean = -np.mean(test_scores, axis=1)
            else:
                train_mean = np.mean(train_scores, axis=1)
                test_mean = np.mean(test_scores, axis=1)

            axes[i, j].plot(train_sizes, train_mean, 'o-', color="r",
                             label="Training")
            axes[i, j].plot(train_sizes, test_mean, 'o-', color="g",
                             label="Cross-validation")
            axes[i, j].set_title(f"{name}: {label}")
            axes[i, j].set_xlabel("Training Samples")
            axes[i, j].set_ylabel(label)
            axes[i, j].legend(loc="best")
            axes[i, j].grid(True)

    plt.tight_layout()
    plt.show()

# Run the function (Uses the 3513 training rows)
plot_comprehensive_learning_curves({"XGBoost": best_xgb, "Bagging RF": bag_rf}, X_train_scaled, y_train)
```



```

In [293]: import seaborn as sns
          import matplotlib.pyplot as plt

          def plot_error_histograms(models, X_test, y_test):
              """
                  Generates histograms of the prediction errors (residuals) for both
                  models.
              """
              fig, axes = plt.subplots(1, len(models), figsize=(15, 6))

              for i, (name, model) in enumerate(models.items()):
                  preds = model.predict(X_test)
                  residuals = y_test - preds

                  sns.histplot(residuals, kde=True, ax=axes[i], color='teal' if
"XGB" in name else 'orange')
                  axes[i].set_title(f"Error Distribution: {name}")
                  axes[i].set_xlabel("Dollar Error (Actual - Predicted)")
                  axes[i].set_ylabel("Frequency")

                  # Add a vertical line at 0 (Perfect Prediction)
                  axes[i].axvline(0, color='red', linestyle='--')

              plt.tight_layout()
              plt.show()

# Run for your models
plot_error_histograms({"XGBoost": best_xgb, "Bagging RF": bag_rf},
X_test_scaled, y_test)

```

